

## **Middleware – Final Technical Requirements for Renesis**

### **Central Function**

- The middleware acts as the **central “nervous system”** of the entire platform.
  - It controls and orchestrates:
    - **AI communication** (Coach Klaus)
    - **Timing and system states**
    - **Feedback cycles**
    - **Dynamic coach role behavior**Discovery\_Detailed\_Proj...
- 

### **Core Technical Requirements**

- **Ultra-low latency and high performance** in communication with the AI (real-time Coach Klaus interaction)
  - **Robust, traceable layer system** for decision logging and explainability
  - **High interface compatibility** with third-party tools:
    - Make.com (workflow automation)
    - LMS editor (content control)
    - Pipedrive CRM (client & lead tracking)Discovery\_Detailed\_Proj...
- 

### **Middleware-Orchestrated Components**

- **Session Memory** (MongoDB-based)
  - **Prompt Engine** (AI control and response logic)
  - **Adaptive Decision Engine (ADE)**
  - **Confidence Score**
  - **Process Intelligent Engine (PIE)**
  - **AI interaction audit logging** (fully AI Act compliant)Discovery\_Detailed\_Proj...
-

## Interfaces & Integration Layer

- Central orchestration of **REST-compliant APIs**
  - All AI-related processes with **D-ID, ElevenLabs, DeepL** are **mediated via the middleware**
  - All personal data **remains within Hetzner (Germany)** – no persistent storage at AWS
  - Planned **GDPR-compliant audit layer** for Aleph Alpha API callsDiscovery\_Detailed\_Proj...
- 

## Security, Monitoring & Transparency

- **AI action traceability** (“Reason Trace”) for every AI decision (trigger, change, justification, timestamp)
  - **Full compliance with the EU AI Act** – logging, explainability, auditability
  - Role-based access control to ensure separation between content, AI logic, and user data
  - **Monitoring dashboards** (admin & trainer) for:
    - Interaction quality
    - Confidence score evolution
    - Adaptive decision historiesDiscovery\_Detailed\_Proj...
- 

## Scalability & Communication Infrastructure

- Middleware handles **asynchronous response pipelines** for Coach Klaus to enable smooth UX
- Designed for **horizontal scaling** (Docker / Kubernetes-ready if needed)
- Clear **communication architecture**:
  - **AWS-based AI services are decoupled**
  - Middleware mediates **only temporary dialog-specific content**
  - No personal data is transferred outside the Hetzner environment

## Session Memory – Technical and Functional Requirements

### Function & Purpose

- The Session Memory serves as the **memory of the platform**, storing contextual user interaction data across modules and over time.
  - It enables:
    - Adaptive responses by Coach Klaus
    - Personalized learning paths
    - Referencing past experiences and decisions
- 

### Data Contents & Structure

- Stores:
    - Progress per module and match day
    - Confidence Scores (per module)
    - Tactics check results
    - User responses to reflection and diagnostics
    - Coach interaction history and role changes
    - Process insights from the PIE module
- 

### Technical Specifications

- **Technology:** MongoDB (NoSQL), fully hosted on Hetzner (Germany)
  - **Structure:** JSON-based entries with timestamp
  - **Data separation:** strict distinction between personal data and technical identifiers
  - **Access control:** only accessible via middleware, trainer dashboards, and admin – *not visible to users or editorial team*
  - **Multi-tenancy:** Separate memory per law firm, optional client-level context
-

## System Integration

- Essential for:
    - **Adaptive Decision Engine** (e.g. coach role switching, learning path adjustments)
    - **Confidence Score** evolution
    - Reminder and follow-up logic
    - Personalized suggestions from Coach Klaus
  - Coach Klaus uses session memory to:
    - Reference earlier conversations and answers
    - Recognize strengths, weaknesses, and emotional phases
    - Recall results from the tactics check
- 

## Data Protection & Compliance

- Stored **exclusively on Hetzner servers** in Germany
- No external sync or export
- Compliant with:
  - **GDPR** (incl. pseudonymization of personal data)
  - **EU AI Act** (traceability and auditability)
- Not visible to:
  - External tools
  - Editorial team
  - CRM or LMS editors

## Adaptive Decision Engine (ADE) – Functional & Technical Requirements

### Core Function

- The ADE is the **emotionally intelligent core** of the platform.
  - It controls:
    - each user's **individual learning journey**
    - the **dynamic role behavior** of Coach Klaus
    - reminders, difficulty levels, repetition logic, bonus unlocksDiscovery\_Detailed\_Proj...
- 

### Input & Output Structure

#### Typical Inputs:

- user\_id, law\_firm\_id, module\_id
- Confidence score, last result, error type
- Session memory context, tactics check results
- Progress state, inactivity duration

#### Typical Outputs:

- coach\_role (e.g., motivating, reflective, challenging)
  - learning\_path\_modifier (e.g., activate repetition, unlock bonus)
  - reminder\_trigger (e.g., time-based prompts)
  - reason\_trace (transparent justification stored in audit log)Discovery\_Detailed\_Proj...
- 

### Rule Logic & Control

#### Rule-based decision engine with editorial access:

- Configurable rule matrix (e.g., via JSON or GUI-based rule editor)
- Examples:
  - $CS < 0.6 \rightarrow \text{Coach} = \text{compassionate}$
  - Repeated errors  $\rightarrow$  Coach explains actively
  - Inactivity  $\rightarrow$  Reminder is triggered

### Trigger points include:

- Completion of a module
- Start of a new match day
- Evaluation of the tactics check
- Inactivity (e.g., after 72 hours)
- Voice input signals like: “I don’t understand that”

---

### Logging & Explainability ("Reason Trace")

- Every decision is **fully logged and explainable** in JSON format.
- Example entry:

json

{

"event": "role\_switch",

"from": "motivating",

"to": "reflective",

"trigger": "confidence\_score\_drop",

"reasoning": "CS of 0.8 → 0.5 after module 3, 2 incorrect answers, long response time",

"timestamp": "2025-06-20T12:34:56"}

- Purpose:
  - AI Act compliance & auditability
  - User-facing transparency
  - Support for quality control & A/B testing

---

### Future Strategy – Hybrid Model

- Initially 100% rule-based – easy to maintain and audit.
- Future: Option to enhance decision logic with machine learning:
  - e.g., decision trees or reinforcement learning
  - Always with a focus on explainability and editorial control

## Prompt Engine – Functional & Technical Requirements

### Purpose & System Role

- The Prompt Engine is the **central control unit for all AI responses** on the platform.
  - It enables:
    - Full control over content, tone, role behavior, and language quality of Coach Klaus
    - Ongoing development and refinement of prompts by the editorial and tech teams
    - Adaptive prompt generation based on user context and learning progressionDiscovery\_Detailed\_Proj...
- 

### Technical Requirements

- Fully **custom-built** – no external prompt service providers
  - Must support **rule-based and version-controlled prompt orchestration**
  - Integrates tightly with:
    - **Session Memory**
    - **Adaptive Decision Engine (ADE)**
    - **Process Intelligent Engine (PIE)**
    - **Confidence Score**
  - Supports:
    - Initial prompts (e.g., onboarding, reflection, coaching)
    - Dynamic prompts (e.g., in response to hesitation, errors, success)
  - Must allow for **real-time performance** (especially for voice-based outputs)
- 

### Monitoring & Analytics

- Full **logging of all prompt–response interactions**
- Dedicated **dashboard** for editorial and technical review:
  - Response quality
  - Latency and reaction time

- Role tracking (e.g., coach role switching over time)
  - Suggestions for prompt adjustments
  - Purpose: quality control, optimization, and transparent auditingDiscovery\_Detailed\_Proj...
- 

### **Security & Auditability**

- All prompts and responses must be:
    - **Documented, versioned, and traceable** (EU AI Act compliance)
    - Logged with full context (user ID, module, timestamp, interaction history)
  - **Edit permissions limited** to specific roles (e.g., editorial, admin)
  - **No external usage** of prompt data for third-party analysis
- 

### **Advanced Functions & Outlook:**

- **A/B testing** of different prompt variations
- Dynamic prompt adaptation via learning analytics
- AI-assisted prompt refinement (e.g., emotional tone adjustment based on confidence signals or user mood)



## Confidence Score – Functional Overview & Requirements

### Purpose & Philosophy

- The Confidence Score measures a user's **subjective sense of certainty** – not their factual accuracy.
- Its primary role is to detect emotional states such as **uncertainty, overload, or hesitation**, and **respond empathetically**.
- It is **not used for blocking or grading**, but to dynamically **adjust the behavior of Coach Klaus** in a supportive wayDiscovery\_Detailed\_Proj...


---

### Calculation Logic (Middleware-Controlled)

- The score is calculated automatically in the middleware based on **weighted behavioral signals**:

Signal Type	Examples	Weight
Error Behavior	Wrong answers, logical errors	high
Demand Behavior	Frequency & type of clarification questions	medium
Usage Signals	Interruptions, repeated actions, pauses	medium
Voice Patterns	Phrases like “I’m not sure...”	high
Speed & Reaction	Hesitation, rushed clicking	low

- Output: **Score between 0% (very uncertain) and 100% (very confident)**
- Continuously updated – e.g. after modules, key interactions, or coach promptsDiscovery\_Detailed\_Proj...

 **Important:** All scoring parameters (weights, thresholds, signal definitions) are still under development.

Therefore, the middleware must support **maximum flexibility and configurability** – including real-time adjustments by the editorial team.

---

### Storage & History

- Score is stored **per user and per module** in the Session Memory
- Visible to **trainer and admin dashboards**, showing trends and history
- **Not visible to users directly** – but expressed through Coach Klaus' behavior

---

## System-Level Impact

### Low Score:

- Coach Klaus becomes more empathetic, explanatory, repetitive
- Reminders may be triggered
- Content complexity is lowered when needed

### High Score:

- Coach Klaus becomes more dynamic, challenging, encouraging
- Bonus content or reflection challenges may be unlocked
- Examples: “You seem confident – want to apply your knowledge?”

 Coach Klaus always stays **supportive and encouraging** – never punitive or limiting  
Discovery\_Detailed\_Proj...

---

## Visualization & Dashboards

- Admin and trainer dashboards show:
  - **Progression curves** per user and firm
  - **Heatmaps** for team sentiment and development
  - **Alerts** for sustained low scores → triggers for learning support

---

## Unique Didactic Feature

- The Confidence Score enables a **new level of emotionally intelligent AI interaction**:
  - Learners feel seen, supported, and understood
  - Learning resistance can be identified without verbal input
  - Coach Klaus becomes a **true companion**, not just a feedback engine

## Process Intelligent Engine (PIE) – Functional & Technical Requirements

### Purpose & Concept

- The PIE is the platform's **real-world process analysis and application engine**.
  - It bridges the gap between learned content and real law firm operations – enabling **practical transfer and reflection**.
  - It is not a modeling tool – but an **AI-supported, empathic analysis companion** for daily workflows.
- 

### Input Types

- **Free-form input** via text or speech  
→ e.g. spontaneous description of a current task flow
  - **File uploads** (e.g., PDF, DOCX, scans)  
→ e.g. internal instructions, protocols, checklists
  - Optional: OCR preprocessing for scanned documents
- 

### Processing via ErxlebenAI (LLaMA 3.3)

- The semantic analysis is performed entirely by **ErxlebenAI (based on LLaMA 3.3)**
  - Prompts are structured via the platform's **custom Prompt Engine**
  - All processing is **GDPR-compliant** and hosted in a secure Hetzner environment
  - The AI identifies:
    - Process steps, role assignments, media transitions
    - Bottlenecks, redundancies, semantic gaps
    - Matches and mismatches with learned content
- 

### Two-Phase Engine Structure

1. **Analysis Module**
  - Extracts, structures, and standardizes process logic
  - Outputs a structured internal JSON representation
2. **Evaluation Module**
  - Compares with module content and best practices

→ Provides a qualitative assessment and concrete, didactically linked suggestions

### Example Output:

json

KopierenBearbeiten

```
{  
  "issues": ["media discontinuity", "lack of automation"],  
  "recommendations": [  
    "Use a digital invoice workflow (see module 4)",  
    "Clarify handover roles and responsibilities"  
  ],  
  "evaluation": "optimizable",  
  "confidence": 92  
}
```

---

### Integration with Coach Klaus

- Coach Klaus is **directly involved** in the process analysis:
  - Actively asks for process descriptions (“Tell me how you currently handle it...”)
  - Explains the evaluation in natural language
  - Links to appropriate modules for further learning
  - Asks follow-up questions if key details are missing
  - Motivates and encourages improvement

---

### Visibility & Dashboards

- Results appear in the **trainer dashboard**
- Multiple processes per firm can be tracked over time
- In future: **internal benchmarking within the same law firm** (no firm-to-firm comparison)

---

## ❌ Distinction from Classic Tools & Documentation Requirement

- PIE is **not** a diagramming or BPMN tool:
  - No technical knowledge or modeling is needed
  - Natural language suffices
  - No specialized software or formats required

## ✅ Nonetheless, a documentable output is essential:

All analyzed processes must be presented in an **exportable and visual form** (e.g. flowchart or structured JSON), suitable for:

- Internal communication
- Coaching sessions
- Team handovers
- The goal is not standardization, but a **practical, easy-to-understand visual reflection** of real-world workflows – flexible, shareable, and aligned with the coaching narrative.

## Learning Management System (LMS)

### System Architecture & Core Setup







- Fully **custom-built LMS** – no third-party platform
  - Fully integrated with:
    - Middleware
    - Session Memory
    - Prompt Engine
    - Dashboards
    - ErxlebenAI (based on LLaMA 3.3)
  - Hosted in a **GDPR-compliant** Hetzner environment (Germany)
  - Scope: **360 match days** = 12 modules × 3 leagues × 10 match days
- 





### Learning Formats & Match Day Structure

#### **Fixed structure per match day: 5+1 phases**

1. Attunement
  2. Knowledge impulse
  3. Processing
  4. Application
  5. Conclusion & motivation
- +1: Resilience impulse (optional)

#### **Key learning formats:**

-  Videos (coach, story, explainer)
-  Quizzes with error classification (input for Confidence Score)
-  Reflection tasks (free text, multiple choice, scaling)
-  Mini simulations (e.g., branching decision trees)
-  Interactive dialogues with Coach Klaus (text, speech, avatar)
-  PIE elements (process reflection via upload or free input)

-  **Process Designer:** Interactive visual builder (e.g., swimlanes, drag-and-drop), exportable
  -  Digital pinboards for team collaboration and idea sharing
  -  “Break Tea” audio: motivational quotes, stadium vibe
  -  League feedback: progress tracking, points, badges, league advancement
- 

## **Tactics Check – Technical Requirements for Renesis**

### 1. **Booking & Scheduling**

- **Booking via LMS-integrated store**
    - “Tactics Check” offered as a bookable product
    - Triggers status change and redirects user to welcome page
  - **Scheduling via calendar tool (e.g. Cal.com)**
    - Full integration into LMS booking flow
    - Data passed to system: user\_id, law\_firm\_id, appointment, booking\_id
    - GDPR-compliant setup required
- 

### 2. **Preparation Phase (Digital)**

- **Pre-check questionnaire (structured form)**
    - Includes scale-based items, checkboxes, free text
    - Data stored per user & law firm in the **Session Memory**
    - Displayed in **trainer dashboard** (with timestamp & status)
  - **Coach Klaus welcome video**
    - Auto-triggered upon completion of the questionnaire
    - Logic-controlled visibility (based on completion)
- 

### 3. **Physical Package Trigger (Optional)**

- **System flag activated 7 days prior to the check**

- Marks user for manual/automated mailing
  - Exportable via CSV/API for logistics provider
  - *Note: LMS does not perform shipping – only triggers export*
- 

#### 4. 🧠 On-Site Execution (Trainer-led)

- **Trainer dashboard (tablet/laptop optimized)**
    - 5–10 scale questions per module + free text
    - Real-time input stored to structured JSON:  
user\_id, modul\_id, score, text, timestamp
  - **Coach Klaus intro video (optional)**
    - Launchable on-site from trainer device
    - Opens the check with motivation and guidance
- 

#### 5. 🎯 Dartboard Animation (Visual Evaluation)

- **Pre-built animation (e.g. SVG or Lottie)**
    - A dart flies into the board **based on score (1–6)**
    - Dynamic position controlled via middleware/frontend
    - Animation is emotional and sports-themed
  - **Optional:** Coach Klaus comment (voice or text) after impact
- 

#### 6. 🏆 Evaluation & League Assignment

- **Automatic league assignment per module**
    - Output: liga = startelf / taktgeber / spielmacher
    - Controlled via middleware (or Adaptive Decision Engine)
    - Impacts: learning path, Coach behavior, content difficulty
- 

#### 7. 🔗 Data Flow & System Integration

- **All check data flows into:**
  - **Session Memory** (contextual baseline)



- **Adaptive Decision Engine** (initial config)
- **Prompt Engine** (for Coach Klaus dialogue references)
- **Trainer Dashboard** (visual feedback, export, alerts)

## **Coach Klaus – Technical Specification & Control Logic**




### **Role in the System**

- Coach Klaus is the **central AI interface** for user interaction throughout the learning journey.
- He acts in **multiple adaptive roles** (e.g., motivating, reflective, explanatory, challenging).
- Responsibilities include:
  - Opening and closing each match day
  - Reflecting on user progress
  - Providing emotional encouragement and clarification
  - Delivering dynamic feedback and reminders

---

## **Core Components & Architecture**

### **1. Multimodal Speech Interface:**

-  **Speech-to-Text** (user input): via ElevenLabs Whisper (or equivalent)
-  **Text-to-Speech** (output): via ElevenLabs API
-  **Avatar rendering**: D-ID HQ Full Body Avatar
- All media orchestrated in real time through the **middleware and prompt engine**

### **2. Prompt Engine Integration:**

- Dynamic prompt generation based on:
  - user\_id, module\_id, match\_day, confidence\_score, coach\_role, and session history
- Prompts control:
  - Content, tone, complexity, media format, and coach behavior

### **3. Adaptive Role Control via ADE:**

- The **Adaptive Decision Engine** determines:

- Role switching
- Coach tonality
- Depth of explanations
- Response frequency
- Based on real-time context and Confidence Score

#### **4. Session Memory Access:**

- Klaus has persistent access to:
    - Previous questions and answers
    - Learning performance and interaction patterns
    - Tactics check results
    - Emotional cues and break history
  - Accessed live through the middleware, no local caching
- 

#### **Interaction Logic**

- **Proactive Triggers:**
    - Start or end of a match day
    - User inactivity or re-entry after absence
    - Milestones like league promotion or PIE success
  - **Reactive Prompts:**
    - User questions via text or voice
    - Coach adjusts behavior and language based on current confidence level
  - **Fallback Logic:**
    - Detects uncertainty or confusion
    - Offers repeat explanations or simplifications
    - Adapts tone and media accordingly
- 

#### **Output & Media Handling**





- **AI-generated real-time responses:**

- Text → TTS → Avatar video (live-rendered)
  - Must render and respond **in less than 2 seconds total latency**
  - **A small number of predefined video responses:**
    - Used for **emotionally significant standard moments** (e.g., start of journey, league promotion, end of module)
    - Stored as fixed assets and referenced via prompt ID
    - Selected by the prompt engine for latency optimization and emotional resonance
- 

## Logging & System Control

- All interactions are fully logged
  - Data is available in:
    - Prompt logs (for optimization and audit)
    - ADE trace log (for role switching decisions)
    - Trainer dashboard (for visibility and analysis)
- 

## System Requirements

-  **End-to-end latency must be < 2 seconds** (including prompt generation, TTS, avatar rendering)
-  Continuous session context access (via middleware)
-  Full logging & traceability (AI Act compliance)
-  Manual trigger option for trainers (e.g., play predefined message)

## Trainer & Admin Dashboards – Technical Requirements

### Purpose & Function

The dashboards serve as the central **monitoring, insight, and control interface** for:

- **Trainers** (learning progress support and coaching)
  - **Admins** (platform and user management)
  - **AI Monitoring Role** (prompt logging, behavior trace, auditability)
- 

### Role-Based Access Logic

Role	Access Scope	Permissions
Trainer	Law firm-specific learning data	Read-only, provide recommendations, no content/system modification
Editorial Team	Content formats & maintenance	No access to user data or logs
Admin (Management)	Full platform access	User, system, content, prompt & audit log control
AI Monitoring	Prompt Engine & model logs	Full AI trace visibility for response analysis & latency tracking

---







### Live Data Sources

All dashboard data is dynamically retrieved via the middleware from:

- **Session Memory** (context, score history)
  - **LMS** (module progress, match day status)
  - **ADE & Confidence Score Engine** (decisions, role logic)
  - **Prompt Engine Logs** (for Coach Klaus tracking)
  - **PIE results** (process evaluations & benchmarks)
  - **Tactics Check results** (initial league assignment, scores)
- 

### Trainer Dashboard – Functional Scope

- Firm- and user-specific overviews

- Drill-down by module, match day, and league
  - Visualizations include:
    -  Tactics Check dashboards
    -  Module & match day completion
    -  Confidence Score timelines & heatmaps
    -  Coach Klaus usage (role, duration, trigger reason)
    -  Inactivity tracking (last login, drop-off point)
    -  PIE usage & process analysis results
  - **Exports:**
    - CSV / PDF
- 

## **🔧 Admin Dashboard – Additional Capabilities**

- System-wide insights:
  - Module and media usage rates
  - Active users per league
  - Booking statuses
  - Drop-off rates & match day conversion
  - Prompt engine stats and rendering performance (Coach Klaus)
  - Badge and point distributions
- Control functions:
  - User and role assignments
  - Visibility overrides (e.g., test modes)
  - Manual content unlocking
  - Module activation (e.g., add-ons, events)
- Audit & compliance features:
  - Prompt & ADE trace export (AI Act compliant)
  - Latency logging (TTS + Avatar)
  - Error tracking & alert logic

## Store & Checkout System – Technical Requirements

### Purpose



- Enables **digital booking and management** of:
    - Modules & course add-ons (e.g., PIE, certifications)
    - Events (e.g., Tactics Check, live sessions)
    - Upgrades (e.g., league transitions, bonus content)
  - Fully integrated into the **LMS and middleware**
  - Checkout must be **GDPR-compliant, secure, and automated**
- 

### Architecture & Components

#### 1. Store Frontend (within LMS):

- Displays available products dynamically based on user role and status
- Filterable by:
  - Module
  - League level
  - Availability (already booked = hidden)
- Supports product visibility logic (e.g., show PIE only to "Playmaker League")

#### 2. Checkout Integration:

- Supported providers:
  -  Stripe (primary)
  -  PayPal (optional)
- Uses redirect or embedded mode
- Confirmation via **Webhook to middleware**

#### 3. Access Activation:

- Upon successful payment:
  - Product is unlocked in the LMS
  - Role and visibility rules are updated
  - Middleware receives all transaction metadata

---

## Product Structure & Metadata

Each store item includes:

- Unique product\_id
- Product type (module / service / download / certificate)
- Visibility logic (public / league-based / personalized)
- Price, booking status, and unlock behavior
- Optional: triggers for content release (e.g., auto-start of match day 1)

---

## Middleware Logic

- After payment, the checkout system calls the middleware via API
- Payload includes:
  - User ID
  - Product ID
  - Payment details
  - Time of transaction
- The middleware:
  - Activates content inside the LMS
  - Updates Session Memory and Adaptive Decision Engine (if applicable)
  - Syncs status with the Trainer & Admin Dashboards

---

## Data Privacy & Compliance

- **No payment data is stored** on the platform
- Only transaction metadata is retained (e.g., Stripe token, webhook ID)
- Fully GDPR-compliant:
  - Uses EU-based payment providers
  - Logs: product booked, timestamp, user ID, receipt reference
  - Invoices handled via middleware or external tool (PDF via email)

---

## Additional Features

- **Promo codes / vouchers** via Stripe API (optional)
- **PDF invoicing** via external service or custom middleware logic
- **Certification unlocking** based on:
  - Module completion
  - Valid payment
  - Admin verification (if needed)

## Multilingual Intelligence – Coach Klaus

(Powered exclusively by ElevenLabs & DeepL)

### Purpose

- Coach Klaus must **automatically detect the user's spoken language** and respond **fluently in that same language**, without any manual language selection.
- This applies to **all supported ElevenLabs languages**, not just German, English, and French.
- The system uses:
  - **ElevenLabs** for both speech recognition (STT) and voice output (TTS)
  - **DeepL** for translation of system prompts where needed

---

## Technical Components

### 1. Language Detection & Transcription (ElevenLabs STT):

- User speaks in any supported language
- ElevenLabs detects:
  - The **spoken language**
  - The **transcribed text**
- Example:

json



KopierenBearbeiten

```
{  
  "language": "it",  
  "text": "Non capisco questo modulo"  
}
```

## 2. Prompt Generation & Language Control:

- Middleware receives both text and detected language
- The Prompt Engine:
  - Constructs a context-aware prompt in the **user's language**
  - If no localized version exists, the German base prompt is translated on-the-fly via **DeepL API**
- The localized prompt is passed to **ErlebenAI (LLaMA 3.3)**
- The model responds **directly in the detected language** (no reverse translation)

## 3. Voice Output via ElevenLabs TTS:

- The AI response is converted to speech using **the correct voice profile per language**
- Each supported language is mapped to a suitable Coach Klaus voice
- Output reflects role, tone, and emotional context

## 4. Avatar Rendering via D-ID:

- TTS audio is synced with D-ID's full-body avatar engine
- The avatar lip-syncs naturally in the user's language – no video pre-rendering required
- Result: Coach Klaus speaks **directly and emotionally in the user's native language**

---

### Middleware-Controlled Interaction Flow

1. User speaks in, e.g., Italian
2. ElevenLabs STT detects "it" and returns transcribed text
3. Middleware passes content + language to the Prompt Engine
4. If needed, DeepL translates the prompt into Italian

5. ErxlebenAI generates the reply in Italian
  6. ElevenLabs converts it to Italian speech
  7. D-ID renders the avatar lip-synced to the Italian audio
  8. User receives a seamless native-language reply
- 

## Technical Requirements

- **ElevenLabs STT & TTS:**
    - Must support all desired languages (DE, EN, FR, IT, ES, PT, PL, NL, etc.)
    - Coach Klaus voice mappings per language (e.g., coach\_voice\_nl)
  - **DeepL Pro API:**
    - For dynamic translation of base prompts if language variant not available
    - Handles all supported DeepL languages automatically
  - **D-ID Full Body Avatar API:**
    - Accepts multilingual audio inputs
    - Renders realistic face & body sync in any spoken language
  - **Middleware must:**
    - Manage detection, translation, prompt routing, and media coordination
    - Track language preference per session (but input language always overrides)
    - Maintain latency **below 2 seconds total** (STT → AI → TTS → avatar)
- 

## Benefits

- Users can **speak freely in their native language**
- Coach Klaus detects and responds **instantly and fluently**
- No settings, toggles, or delays
- Fully scalable to new languages as ElevenLabs expands voice support

